

Received 17 March 2025, accepted 21 April 2025, date of publication 25 April 2025, date of current version 5 May 2025.

Digital Object Identifier 10.1109/ACCESS.2025.3564530

RESEARCH ARTICLE

Multimodal and Temporal Graph Fusion Framework for Advanced Phishing Website Detection

S. KAVYA¹ AND D. SUMATHI¹

School of Computer Science Engineering and Information Systems, Vellore Institute of Technology, Vellore 632014, India

Corresponding author: D. Sumathi (dsumathi@vit.ac.in)

This work was supported by Vellore Institute of Technology, Vellore, India.

ABSTRACT Phishing attacks are among the persistent threats that are dynamically evolving and demand advanced detection mechanisms to counter more sophisticated techniques. Traditional detection approaches are usually based on single-modal features or static analysis, failing to capture the complex, multi-faceted nature of phishing websites and their dynamic behaviors. Thus, we present a robust Multi-Modal and Temporal Graph Fusion Framework integrating advanced learning paradigms that enhance accuracy and adaptability in phishing detection. Our work proposes four brand-new methods: Multi-Modal Hypergraph Fusion Network (MM-HFN), Temporal Graph Neural Network with Attention (TGNN-Att), Federated Graph Contrastive Learning Network (FGCL-Net), and Multi-Modal Temporal Hypergraph Fusion Network (MMTHF-Net). MM-HFN leverages hypergraphs to capture complex, high-order relationships at textual levels (BERT) and graph-based features versus visual ones (CNNs) for an accuracy in the 95-97% range. TGNN-Att addresses temporal variations in phishing behavior by using attention-enhanced temporal graph networks and LSTMs, providing dynamic detection with 94-96% accuracy. FGCL-Net ensures privacy-preserving learning across decentralized datasets through federated contrastive learning, achieving 93-95% accuracy while safeguarding data privacy. Finally, MMTHF-Net fuses multi-modal and temporal features into a dynamic hypergraph framework, achieving state-of-the-art accuracy of 96-98% with an F1-score of 0.97. These approaches together allow for exact, real-time phishing detection by capturing static and temporal behaviors, high-order relationships, and cross-modal features. The framework proposed demonstrates significant improvements compared to the state of the art, eliminating the shortcomings of single-modality and static analysis while offering scalability, privacy, and adaptability levels.

INDEX TERMS Phishing detection, multi-modal learning, hypergraph networks, temporal analysis, federated learning, scenarios.

I. INTRODUCTION

Phishing websites have turned out to be one of the most pervasive cyber threats, using vulnerable users to steal personal information, financial assets, and organizational security. With the continually increasing complexity of phishing attacks and very fast developments in evasion mechanisms, traditional detection systems can no longer suffice. Traditional approaches [1], [2], [3] greatly rely on single-modal inputs and often just resort to text-based

heuristics, static URL features, and even the structure of a particular webpage. Such methods cannot capture the overall multi-dimensional nature of phishing attacks. Most approaches tend to miss phishing website behavior, which most of the time is dynamic in terms of how the structure, content, or hosting domains of phishing pages are adjusted to evade detection. Such limitations point to the pressing need for advanced detection mechanisms that are capable of analyzing complex relationships, multi-modal features, and temporal variations simultaneously. This paper addresses these challenges [4], [5], [6] by introducing a comprehensive Multi-Modal and Temporal Graph Fusion Framework for

The associate editor coordinating the review of this manuscript and approving it for publication was Tao Huang¹.

phishing detection, relying on cutting-edge techniques in graph-based learning, multi-modal feature integration, temporal analysis, and privacy-preserving federated learning. These proposed methods bring together the textual (BERT-based), visual (CNN-extracted), and graph-based representations toward a holistic understanding of phishing webpages. MM-HFN models complex, high-order relationships, whereas the Temporal Graph Neural Network with Attention, or TGNN-Att, captures the evolving patterns of phishing behavior across the sets of temporal instances. For federated graph contrastive learning, it is scalable with privacy support through FGCL-Net, enabling decentralized learning under data confidentiality. Finally, with the integration of all modalities into a dynamic temporal hypergraph structure, MMTHF-Net goes on to deliver state-of-the-art performance. These methods open up new ways toward the detection and mitigation of complex phishing attacks with unprecedented precision and flexibility. This work thus fills an important gap that existed in the existing methods: It introduces a robust fusion of multi-modal features, hypergraphs, and temporal dynamics. The experimental results prove the better performance of the proposed models, achieving an accuracy up to 98% and a high F1-score of 0.97. Federated learning is also integrated, which ensures privacy compliance, so the framework is suitable for deployment in diverse and sensitive environments. This research sets a new benchmark for phishing website detection, offering a scalable, adaptable, and privacy-preserving solution to mitigate modern phishing threats.

A. MOTIVATION AND CONTRIBUTION

Rapid phishing site evolutions, and thus evasion over the traditional detection mechanism forms the main critical challenge in the cyber domain. Approaches that have been developed or are based on either textual analysis or static graph-based approaches are insufficiently sensitive to begin with to deal with those attacks that contain misleading visuals, dynamic content, as well as even changes of time. Furthermore, static detection models do not recognize the evolving patterns of phishing because websites frequently change links, domain structures, or server behavior to avoid detection tools. With the increasing requirement for thorough detection methods, it becomes essential to integrate multi-modal features like text, visuals, and behavior with high-order relationships and temporal variations.

Another challenge is the privacy of decentralized datasets in real-world environments where raw data sharing for model training might violate confidentiality agreements or breach data protection regulations. The authors propose a unified Multi-Modal and Temporal Graph Fusion Framework that addresses the above challenges by four innovative methods: MM-HFN, TGNN-Att, FGCL-Net, and MMTHF-Net. These approaches leverage multi-modal learning for fusion of textual, visual, and graph-based representations, high-order modeling via hypergraph networks, and temporal analysis to

capture time-evolving phishing patterns. Notably, federated graph contrastive learning can enable decentralized, privacy-preserving training without transmitting raw data samples. The contributions of this work are: (1) the integration of multi-modal and temporal data into hypergraph-based models enhances detection precision; (2) the use of attention mechanisms improves the identification of critical changes in evolving phishing behaviors; and (3) the federated learning approach ensures privacy compliance while maintaining high accuracy across diverse datasets. Altogether, these innovations achieve the best results in state-of-the-art phishing detection with higher accuracy, adaptability, and privacy preservation.

ABBREVIATIONS

Acronyms

LLM
GAN
CNN
LSTM
ML
DL
GNN
ELM
URL
TLS
RNN
NLP
HTML
SMOTE

RF
MOE
GRU
PDGAN

TEGDetecter
Eth-PSD
DEPHIDES

EWDPs

PEAE-GNN

PhishSim
SmartiPhish

ResMLP
MDepthNet
SMOTE-Tomek

FGCL-Net

MM-HFN

Full Form

Large Language Model.
Generative Adversarial Network.
Convolutional Neural Network.
Long Short-Term Memory.
Machine Learning.
Deep Learning.
Graph Neural Network.
Extreme Learning Machine.
Uniform Resource Locator.
Transport Layer Security.
Recurrent Neural Network.
Natural Language Processing.
HyperText Markup Language.
Synthetic Minority Over-sampling Technique.
Random Forest.
Multi-Objective Evolution.
Gated Recurrent Unit.
Phishing Detection with Generative Adversarial Network.
Transaction Evolving Graph Detector.
Ethereum Phishing Scam Detection.
Deep Learning-Based Phishing Detection System.
Early Warning and Detection on Phishing Scams.
Phishing Detection via Augmented Ego-Graph GNN.
Phishing Simulation Tool.
Intelligent Anti-Phishing Reinforcement Learning Solution.
Residual Multi-Layer Perceptron.
Multi-Depth Network.
SMOTE Combined with Tomek Link Method.
Federated Graph Contrastive Learning Network.
Multi-Modal Hypergraph Fusion Network.

TGNN-Att	Temporal Graph Neural Network with Attention.
MMTHF-Net	Multi-Modal Temporal Hypergraph Fusion Network.
Phish-Armour	Phishing Detection Using Deep RNN.

II. REVIEW INTO STUDIES RELATED TO PHISHING ANALYSIS

Phishing detection has emerged as one of the most pressing issues in cybersecurity, mainly due to sophisticated attack strategies and their fast-evolving nature. Over the last two decades, there have been tremendous improvements to detect phishing websites, emails, and blockchain-based scams through various computational models, ranging from rule-based systems to state-of-the-art machine learning and deep learning methods. Recent developments in large language models (LLMs) have brought about the possibility of implementing new detection systems, where Koide et al. [1] have employed the use of LLMs on textual and code-based feature analysis in phishing sites to be able to improve their performance in accuracy and cognition-driven detection. However, not all the research scopes point to the textual analysis kind.

Zieni et al. [2] presented a comprehensive survey on phishing website detection, categorizing techniques into blacklist-based systems, machine learning approaches, and visual similarity methods. Their work revealed the limitations of blacklist methods in handling zero-day phishing attacks and emphasized the increasing role of machine learning models for dynamic phishing detection. The detection of phishing activities on blockchain platforms, such as Ethereum, has emerged as a critical domain.

Luo et al. [3] presented a phishing account detection model using random walks for identifying malicious Ethereum accounts. The resilience of the detection models was tested through adversarial hiding approaches by Wen et al. [4], who demonstrated how the perturbation methods may mask the phishing features, bringing about new challenges in phishing detection using the Ethereum-based model.

Sahingoz et al. [5] tackled phishing website detection using deep learning techniques, introducing DEPHIDES, which combines various detection algorithms to improve accuracy. Their study highlighted the significance of end-to-end learning in phishing detection, which was further corroborated by Wu et al. [6], who implemented network embedding strategies to identify phishing scams in blockchain networks, achieving improved performance compared to conventional systems. Iteratively, as shown in table 1, Kabla et al. [7] proposed Eth-PSD, a machine learning-based phishing scam detection model. Using feature engineering, it extracted blockchain-specific attributes. One of the common challenges in phishing detection is the evasion capability of phishing websites which uses cloaking techniques.

Li et al. [8] conducted systematic review of these techniques as to how phishing websites employ anti-detection

mechanisms in avoiding security tools. In addressing this, adversarial learning approaches have surfaced as reported by Al-Ahmadi et al. [9] which utilized which utilized the application of Generative Adversarial in conjunction with CNNs and LSTMs for phishing detection. This presents a very strong advantage with regard to generating adversarial samples in the development of a robust model. Similarly, Zhu et al. [10] proposed a multi-objective evolutionary optimization algorithm in combination with a Random Forest classifier to improve the detection accuracy through feature selection. Multimodal learning has also been applied in phishing detection, such as in the work of Duy et al. [11], where adversarial sample resistance was explored for multimodal models by combining textual, visual, and structural features to improve robustness. Phishing email detection is another important area of phishing attacks, which has been widely researched using NLP techniques. Salloum et al. [12] discussed the systematic literature review on phishing email detection based on NLP-based feature extraction and machine learning classification techniques. Feature-free tools, like PhishSim proposed by Purwanto et al. [13], eliminate manual feature engineering to provide efficient phishing detection using end-to-end models. Graph-based techniques are further discussed, as proposed by Kang and Buu [14], in anomaly detection for phishing scams in cryptocurrency transactions using disentangled prototypical autoencoders. Their work highlighted the need for graph anomaly detection in the discovery of hidden relationships in transactional networks. NLP is also used in phishing email creation and detection, as Heiding et al. [15] reported, where LLMs were used in developing phishing emails while simultaneously detecting them as an example of the two-sided nature of AI breakthroughs. Feature-based phishing URL detection is at the center of most of the mechanisms in the detection system. Remya et al. [16] suggested ResMLP approach for the phishing URL classification with optimized pipeline processing with high accuracy. Alsubaei et al. [17] proposed the hybrid deep learning framework along with improvement in combination of ResNeXt and GRUs for features URL, content, and visual. In contrast, Asiri et al. [18] did a survey on intelligent detection designs for HTML-based phishing URLs and emphasized the need for machine learning and NLP in the current phishing detection systems. Enterprise-specific solutions were investigated by He et al. [19], who proposed a double-layered system for the detection of internal threats using deep learning for identifying phishing attacks within the organizational networks. In the malicious URL detection domain, Rafsanjani et al. [20] developed a feature evaluation framework that leverages priority coefficients in order to improve classification models. Ensemble methods are also promising; for example, Wei and Sekiya [21] showed that ensemble machine learning methods can be sufficient for phishing detection. He et al. [22] has enhanced the URL detection technique further with the feature extension technique. Meanwhile, Zheng et al. [23] used the dynamic graph classification approach based on the TEGDetector

model for detecting the evolving phishing behavior. The improvement in the phishing detection mechanism is done based on the GNN model proposed by Huang et al. [24] with PEAE-GNN. Real-time phishing detection systems, for instance, EWDPs as presented by Xu et al. [25], detect phishing scams on the Ethereum platform in real-time and warn users of engagement before it happens. Machine learning approaches have been used with adaptations for adversarial attacks, as can be seen in the study done by Sudar et al. [26], who implemented adversarial training defense mechanisms for phishing attacks. Similarly, Kumar et al. [27] focused on the detection of phishing from encrypted TLS traffic that again brought the need for feature extraction in complex network environments into limelight.

The real-world applications of machine learning-based detection were presented by Bezerra et al. [28] and Jha et al. [29] who have used feature-based and feature-free methods for the detection of phishing websites. Giri and Banerjee [30] further enriched ensemble learning methods by proposing an optimal greedy stacking model for phishing website detection. In email phishing detection, Somesha et al. [31] proposed DeepEPHishNet using word embeddings and autoencoders to boost the accuracy. Hong et al. [32] have taken care of script-based attacks as well as sophisticated adversarial methods with the help of client-side detection mechanisms using machine learning frameworks.

The approaches such as continual learning approaches, for instance, life-long phishing detection by Ejaz et al. [33] are adaptive to dynamic environments. Color-based phishing detection models like Phish-Sight by Pandey and Mishra [34] utilized visual deception techniques using the dominant colors of web pages. Further improvements in ensemble techniques are shown by Pandey et al. [35], where the analysis of phishing websites is carried out on intelligent grounds. The most modern deep learning models, with the GAN-CNN-LSTM framework proposed by Albahadili et al. [36] being one of them, have robust detection capabilities combined with swarm intelligence optimization. Feature selection techniques were optimized by diversity-oriented search algorithms as adopted by Bacanin et al. [37]. Alowaimier and Dahiya demonstrated performance enhancement of deep learning techniques by optimizing CNN architectures for phishing detection purposes [38]. In addition to the above techniques, Reyes-Dorta et al. and Chanis and Arampatzis have researched URL detection by making use of machine learning techniques in combination with stylometric features and classifier stacking to enhance the detection rate of phishing emails [39], [40]. Ponnusamy and Dhandayudam [41] further introduced novel feature selection and ensemble techniques, using bagging learning and siamese networks for phishing detection. Remmide et al. [42] employed siamese networks for phishing detection.

Blockchain-based adversarial detection, such as Yan et al. [43], applied domain adaptation to identify phishing behaviors across different chains. Reinforcement learning

approaches, such as SmartiPhish by Ariyadasa et al. [47], present intelligent anti-phishing strategies, while integrated deep learning methodologies like MDepthNet proposed by Yamarthy and Koteswararao [44] further enhance detection accuracy. Recent developments, such as permutation importance-based deep learning framework proposed by Zaimi et al. [45], focuses phishing URL detection with recent optimisation of feature techniques. Kondaiah et al. [46] focus to detect encrypted traffic and NLP features-based dual layer CNN and GRU are presented in Birthriya et al. [48]. Finally, deep recurrent neural network for phishing detection attacks is proposed as Phish-armour by Dhanavanthini and Chakkravarthy [50]. In conclusion, this review highlights the diverse and changing approaches to phishing detection, including large language models, machine learning, deep learning, graph-based techniques, and blockchain-specific solutions. Although much progress has been made, the evolution of phishing attacks calls for further research into real-time, robust, and adaptable detection mechanisms.

III. PROPOSED MODEL DESIGN

To overcome the problems related to low efficiency and high complexity that are present in current methodologies, this work will involve the discussion of the efficient design of a Multimodal and Temporal Graph Fusion Framework for Advanced Phishing Website Detection Operations. Initially, as per figure 1, The Multi-Modal Hypergraph Fusion Network (MM-HFN) and Temporal Graph Neural Network with Attention (TGNN-Att) are two advanced models which were designed to overcome the constraints of existing phishing detection frameworks by incorporating multi-modal and temporal information sets. The MM-HFN mainly focuses on the static high-order relationships over different modalities whereas the TGNN-Att further enhances the detection through capturing evolving behaviors over the timestamp with dynamic attention mechanisms. These models together form a strong system, capitalizing on the strengths of multi-modal features and temporal dynamics to produce state-of-the-art detection levels. The MM-HFN models phishing detection as a hypergraph classification problem by capturing high-order relationships among webpage components (nodes such as images, links, and metadata) using hyperedges in the process. For a set of nodes V and hyperedges E via equations 1 and 2,

$$V = \{v^1, v^2, \dots, v_n\} \quad (1)$$

$$E = \{e^1, e^2, \dots, e_m\} \quad (2)$$

The hypergraph structure is presented as $G = (V, E, H)$, where H is the incidence matrix for this process. Node features are initialized with multi-modal embeddings, and in the process, combine the BERT-extracted textual features X_t with CNN-extracted visual features X_v . The node feature matrix X is thus defined via equation (3):

$$X = \alpha X_t + \beta X_v \quad (3)$$

TABLE 1. Comparative analysis of existing methods.

Reference	Method Used	Findings	Results	Limitations
[1]	ChatPhishDetector (LLMs)	Detect phishing websites using large language models	Achieved high accuracy through semantic analysis.	Limited scalability to rapidly evolving phishing strategies.
[2]	Survey on Phishing Detection	Comparative analysis of phishing detection techniques	ML methods outperformed blacklists for zero-day attacks.	High computational cost for real-time ML applications.
[3]	Network Embedding for Ethereum	Detect phishing accounts on Ethereum using embedding.	Successfully identified phishing nodes on Ethereum.	Vulnerable to adversarial graph perturbations.
[4]	Adversarial Hiding Approach	Evaluate adversarial hiding against phishing detection.	Highlighted weaknesses in Ethereum phishing detection.	Requires constant model retraining for robustness.
[5]	DEPHIDES (Deep Learning)	Deep learning-based phishing detection framework.	Improved accuracy with classification algorithms.	Overfitting risk on small datasets.
[6]	Network Embedding for Phishing	Use network embedding to detect phishing scams.	Achieved better performance in blockchain phishing.	Struggles with dynamic blockchain environments.
[7]	Eth-PSD (Machine Learning)	ML-based phishing scam detection on Ethereum.	Detected blockchain phishing effectively.	Scalability for real-time Ethereum monitoring.
[8]	Cloaking Technique Review	Systematic analysis of cloaking methods for phishing.	Classified evasion techniques for detection systems.	Lack of real-time implementation.
[9]	PDGAN (GAN + LSTM)	GAN-based phishing detection with deep learning.	Improved adversarial robustness and accuracy.	Computationally intensive.
[10]	MOE/RF (Random Forest)	Optimize phishing detection with multi-objective methods.	Improved detection via feature selection optimization.	Limited to structured datasets.
[11]	Multimodal Adversarial Defense	Study adversarial resistance for multimodal models.	Enhanced robustness against adversarial attacks.	High latency in multimodal fusion.
[12]	NLP-Based Phishing Email Detection	Review of NLP techniques for phishing emails.	Highlighted effectiveness of semantic-based detection.	Lack of evaluation on real-world data samples.
[13]	PhishSim (Feature-Free Tool)	Feature-free phishing detection framework.	Reduced need for manual feature extraction.	Limited to static phishing webpages.
[14]	Graph Anomaly Detection	Autoencoder-based phishing detection on blockchains.	Detected anomalies in cryptocurrency transactions.	Requires large labeled graph data samples.
[15]	LLMs for Phishing Emails	Generate and detect phishing emails using LLMs.	LLMs can generate deceptive and realistic emails.	Dual-edge challenge: creation of phishing content.
[16]	ResMLP-Based URL Detection	Detect phishing URLs using residual MLP pipelines.	Achieved significant improvement in URL classification.	Limited generalizability to non-URL attacks.
[17]	Hybrid Deep Learning Framework	Combine CNN and GRU for phishing detection.	Improved precision and recall in detecting phishing.	Training is resource-intensive.
[18]	Intelligent HTML URL Detection	Survey of HTML-based URL detection techniques.	Emphasized NLP and intelligent ML methods.	No practical deployment solutions.
[19]	Double-Layer Detection	Detect internal phishing threats in enterprise systems.	Detected insider phishing attacks using deep learning.	High computational requirements for real-time systems.
[20]	Feature Evaluation Framework	Prioritize features for malicious URL detection.	Improved detection using feature optimization.	Dependency on extensive feature engineering.
[21]	Ensemble Learning for Phishing	Test ensemble learning for phishing website detection.	Ensemble methods achieved competitive accuracy.	Not suitable for lightweight detection.
[22]	Feature Extension for URL Detection	Enhance phishing detection using feature extension.	Enhanced detection accuracy through feature expansion.	Ineffective on unstructured data samples.
[23]	TEGDetecter (Dynamic Graph)	Detect evolving phishing behaviors using graph models.	Adapted to dynamic phishing changes over temporal instance sets.	Requires continuous monitoring and updates.
[24]	PEAE-GNN for Ethereum	Use GNN for phishing detection on Ethereum.	Achieved high accuracy for graph-based blockchain data samples.	Limited interpretability of GNNs.
[25]	EWDPs (Early Warning)	Early detection of Ethereum phishing scams.	Provided early phishing warnings effectively.	Scalability remains a challenge for large networks.

TABLE 1. (Continued.) Comparative analysis of existing methods.

Reference	Method Used	Findings	Results	Limitations
[26]	Adversarial Phishing Defense	Detect adversarial phishing attacks via ML.	Enhanced detection against adversarial attacks.	Limited adversarial testing on real-world data samples.
[27]	ML for TLS Traffic	Detect phishing from TLS encrypted traffic.	Identified phishing URLs effectively using ML.	Limited detection capability on non-TLS traffic.
[28]	ML-Based Phishing Detection	Machine learning net for website detection.	Improved detection accuracy through ML features.	Requires large labeled datasets.
[29]	Intelligent Phishing Detection	Machine learning-based phishing detection.	Highlighted ML models' ability to classify phishing.	Limited to static datasets.
[30]	Greedy Stacking for Phishing	Use ensemble stacking for phishing detection.	Achieved optimized performance using stacking methods.	High training timestamp for large models.
[31]	DeepEPishNet	Word embedding and deep learning for email detection.	Improved email phishing detection using embeddings.	Computational complexity for deep models.
[32]	ML for Script-Based Attacks	Detect phishing and script-based attacks.	Client-side security improved using ML frameworks.	Limited evaluation on adversarial scripts.
[33]	Life-Long Detection (Continual)	Continual learning for phishing detection.	Improved adaptability to evolving phishing strategies.	Requires constant retraining.
[34]	Phish-Sight (Color-Based)	Use dominant colors on web pages for phishing detection.	Improved visual phishing detection using colors.	Ineffective on text-based phishing.
[35]	ML Ensemble for Phishing Sites	Combine ensemble ML for phishing website detection.	Improved accuracy using ensemble methods.	Computationally expensive for large-scale data samples.
[36]	GAN-CNN-LSTM Framework	Use GANs with deep learning for phishing detection.	Enhanced detection and robustness.	High resource consumption during training.
[37]	Feature Selection + ELM	Optimize feature selection for phishing detection.	Improved detection with fewer features.	Ineffective on unstructured phishing data samples.
[38]	Improved DL for Phishing	Optimize deep learning for phishing website detection.	Achieved competitive performance with DL enhancements.	Limited real-world testing.
[39]	ML-Based Malicious URL Detection	Detect malicious URLs using machine learning.	Successfully identified phishing URLs.	Lack of multimodal analysis.
[40]	Stylometric Email Detection	Use stylometric features for phishing email detection.	Enhanced email detection with classifier stacking.	Limited to stylometry-based features.
[41]	Optimized Bagging Learning	Use bagging ensemble with feature selection.	Improved performance for URL phishing detection.	Not suitable for real-time applications.
[42]	Siamese Deep Learning	Authorship verification for phishing emails.	Robust detection of email-based phishing.	Ineffective for websites and blockchain phishing.
[43]	Adversarial Domain Adaptation	Use domain adaptation for blockchain phishing.	Improved adaptability across blockchain networks.	Requires labeled domain data for training.
[44]	MDepthNet (Integrated DL)	Deep learning framework for phishing attack detection.	Achieved state-of-the-art performance.	Limited feature interpretability.
[45]	Deep Learning with SMOTE	Deep learning with SMOTE-Tomek for phishing URLs.	Balanced datasets improved detection performance.	Requires preprocessing for imbalanced datasets.
[46]	Ensemble Learning for TLS URLs	Detect phishing URLs in encrypted traffic.	Improved detection accuracy for encrypted URLs.	High computation for encrypted traffic analysis.
[47]	SmartiPhish (Reinforcement)	Reinforcement learning for spoofed website detection.	Enhanced intelligent anti-phishing strategies.	Limited scalability in large-scale detection.
[48]	Dual-Layer CNN + GRU	CNN-GRU hybrid for phishing detection.	Improved accuracy using attention and NLP features.	Computationally expensive for real-time systems.
[49]	Ransomware Detection Service	Machine learning-based ransomware detection.	Identified ransomware attacks effectively.	Not specialized for phishing attacks.
[50]	Phish-Armour (RNN-Based)	Deep RNNs for phishing detection.	Improved detection accuracy using recurrent models.	High latency for real-time systems.

where α and β are modal scaling factors. Hypergraph convolution is employed to spread features across the nodes and hyperedges within this process. Via equation (4), the model defines the operation of convolution:

$$Z = D_V^{-\frac{1}{2}} H W D_e^{-\frac{1}{2}} H^T D_V^{-\frac{1}{2}} X \quad (4)$$

where D_V and D_e are diagonal matrices representing node and hyperedge degrees, respectively, and W is the learnable weight matrix for this process. The output feature matrix Z preserves high-order relationships between nodes. The classification probabilities for nodes are obtained by passing Z through a softmax layer via equation (5):

$$P_i = \text{softmax}(W_o Z_i + b_o) \quad (5)$$

where W_o and b_o are learnable parameters, and P_i represents the phishing probability for node 'i'. Iteratively, the TGNN-Att extends the detection capability of the model towards dynamic phishing behaviors. This is done by incorporating temporal snapshots. For each timestamp t , a temporal graph $G_t = (V_t, E_t)$ is defined by the process, where nodes V_t denote webpage components, and edges E_t capture temporal relationships such as transitions or redirects. Node features at timestamp t , X_t , evolve over timestamps and are fed into a temporal graph neural network for processing purposes. The temporal propagation is defined using self-attention mechanisms. For each node i , the attention weight between node i and its neighbor j is computed via equations (6), (7), and (8):

$$\alpha_{ij}^t = \frac{\exp(q_i^T k_j)}{\sum_{k \in N(i)} \exp(q_i^T k_k)} \quad (6)$$

$$q_i = W_q X_{it} \quad (7)$$

$$k_j = W_k X_{jt} \quad (8)$$

where W_q and W_k are learnable weight matrices, and $N(i)$ denotes the neighbors of node i in the process. The updated node representation h_i^{t+1} is obtained via equation (9):

$$h_i^{t+1} = \sigma \left(\sum_{j \in N(i)} \alpha_{ij}^t W_v X_{jt} \right) \quad (9)$$

where σ is a non-linear activation function and W_v is a learnable weight matrix for value transformations. To capture long-term dependencies across the temporal instance sets, updated node representations H_t are put through an LSTM process represented at the high level via equation (10):

$$H_t = \text{LSTM}(H_{t+1}, h_t) \quad (10)$$

At a timestamp of t , H_t signifies the hidden state that embeds both short-term and long-term behaviors of phishing nodes. The overall phishing probability for each node is determined via equation (11):

$$P_t = \text{softmax}(W_o H_t + b_o) \quad (11)$$

This way, the temporal self-attention with LSTM layers ensures TGNN-Att correctly highlights the important

changes in time - links, redirects, or changes in resource structure. Thus, the choice of MM-HFN and TGNN-Att is correct since both have complementary strengths. This work, on one hand, excels at modeling static multi-modal relationships using hypergraphs, and TGNN-Att improves the temporal adaptability by capturing evolving phishing behaviors. Thus, together, these techniques form a comprehensive system that attains high detection accuracy with superior adaptability levels. This is done by Exploiting multi-modal embeddings, along with high-order graph convolutions, and temporal self-attention processes. Due to which these models significantly outperform traditional detection methods and provide robust detection against sophisticated phishing techniques. Iteratively, Next, As depicted in figure 2, Intelligently, the Federated Graph Contrastive Learning Network (FGCL-Net) and the Multi-Modal Temporal Hypergraph Fusion Network (MMTHF-Net) are designed to improve phishing website detection efficiency levels. This is done by privacy-preserving feature learning and advanced temporal multi-modal analysis. FGCL-Net focuses on learning high-quality representations in a decentralized setting without sharing raw data, while MMTHF-Net integrates multi-modal features and dynamic relationships using hypergraphs to capture evolving phishing behaviors over temporal instance sets. These methods complement earlier models by incorporating privacy preservation and temporal adaptability while modeling high-order relationships. The FGCL-Net exploits the principles of graph contrastive learning (GCL) and federated learning (FL) to learn generalized phishing detection models across decentralized datasets. Let $G_i = (V_i, E_i, X_i)$ represent a local graph from the i -th client, where V_i is the set of nodes, E_i are edges, and X_i represents node features. GCL generates two Augmented views, G_{i1} and G_{i2} , through structural perturbation, such as node masking or edge dropping, and temporal augmentations. The contrastive loss LGCL is defined via equation 12,

$$LGCL = -\frac{1}{|V|} \sum_{i \in V} \log \left(\frac{\exp(\text{sim}(z_{i1}, z_{i2})/\tau)}{\sum_{j \in V} \exp(\text{sim}(z_{i1}, z_{j2})/\tau)} \right) \quad (12)$$

where z_{i1} and z_{i2} are node embeddings from the two graph views, $\text{sim}(\cdot)$ is cosine similarity, and τ is a temperature scaling factor.




FIGURE 1. Model architecture of the proposed analysis process.

Federated learning aggregates local updates to form a global model without accessing raw data samples. Each local client k optimizes its model θ_k using the loss function:

$$L_k = LGCL + \lambda \|\theta_k\|^2 \quad (13)$$

where λ is the regularization factor. The global model θ is updated using the Federated Averaging (FedAvg) algorithm:

$$\theta = \frac{1}{K} \sum_{k=1}^K \theta_k \quad (14)$$

where K is the total number of clients. The global embeddings Z obtained from FGCL-Net are further classified using a softmax function:

$$P_i = \text{softmax}(W_o z_i + b_o) \quad (15)$$

where P_i is the phishing probability for node i .

The MMTHF-Net builds upon hypergraph structures to fuse multi-modal (textual, visual) and temporal features into a unified framework. Given a temporal hypergraph $H_t = (V_t, E_t, X_t)$ at timestamp t , node features include textual embeddings X_T , visual features X_V , and temporal embeddings X_{temp} . The integrated node feature matrix is represented as:

$$X_t = \gamma_1 X_T + \gamma_2 X_V + \gamma_3 X_{temp} \quad (16)$$

where $\gamma_1, \gamma_2, \gamma_3$ are learnable weighting factors. Temporal hypergraph convolution propagates node features across hyperedges with time-awareness:

$$H_t = D_V^{-\frac{1}{2}} H W_t D_e^{-1} H^T D_V^{-\frac{1}{2}} X_t \quad (17)$$

where W_t is a time-dependent learnable weight matrix, D_V and D_e represent node and hyperedge degrees, respectively. Temporal relationships are further refined using attention mechanisms:

$$\alpha_{ij}^t = \frac{\exp(q_i^T k_j / d)}{\sum_{k \in N(i)} \exp(q_i^T k_k / d)} \quad (18)$$

$$q_i = W_q h_i \quad (19)$$

$$k_j = W_k h_j \quad (20)$$

The updated temporal embeddings are obtained as:

$$h_i^{(t+1)} = \sigma \left(\sum_{j \in N(i)} \alpha_{ij}^t W_v h_j \right) \quad (21)$$

where σ is a non-linear activation function. Finally, phishing probabilities for nodes are computed as:

$$P_t = \text{softmax}(W_o H_T + b_o) \quad (22)$$

The MMTHF-Net integrates multi-modal features, temporal relationships, and high-order node interactions to deliver robust phishing detection. Its ability to model temporal dynamics and fuse diversified features makes it ideal in detecting sophisticated phishing techniques which evolve over temporal instance sets. FGCL-Net and MMTHF-Net




FIGURE 2. Overall flow of the proposed analysis process.

complement each other well in handling the challenges presented by decentralized data privacy and temporal feature fusion. FGCL-Net guarantees privacy-preserving learning across distributed datasets while making it scalable and conforming to data regulations. MMTHF-Net integrates multi-modal and temporal information through hypergraph learning, and achieves state-of-the-art accuracy. Together, they form a comprehensive and adaptive system for phishing detection. The final output of the combined framework represents precise node-level and graph-level phishing probabilities, ensuring both privacy preservation and temporal adaptability levels. All variables consistent with Equations (3) - (22), including scaling factors (α, β, γ) and contrastive loss temperature parameter (τ), were fully optimized by grid search to ensure uniformity among them. The ranges incorporated values between 0.1 and 1.0, changing by 0.1 to tune α and β , while γ values for feature fusion (textual, visual, and temporal components) were tuned between 0.2 and 0.5 to balance multimodal contributions effectively. The resultant cross-validation results ensured that the values of γ were optimal at 0.4 (textual), 0.3 (visual), and 0.3 (temporal), ensuring a balance in the contribution of all modalities and not being skewed by any single modality during detection. Similarly, variation between 0.3 and 1.0 was done on τ of contrastive learning, the best performing setting resulting from 0.5 for generating separation to appear against phishing-innocent samples but with not too close a collapse.

The analysis of the computational complexity of the two networks MMTHF-Net and FGCL-Net was done with scaling up the dataset size. With hypergraph convolution, MMTHF-Net obtains an overall complexity in $\mathcal{O}(|V| + |E|)$ per layer, in which $|V|$ is the number of nodes (webpage components) and $|E|$ hyperedges (multi-modal interactions) in this process. For instances of increasing sizes of datasets, say from 10,000 to as high as 50,000 nodes, empirical runtime analysis on this indicated that training times increased from 220 seconds to 510 seconds per epoch but demonstrated clear trends toward linear scalability. The FGCL-Net aggregates messages from federated averaging executed over 10 decentralized clients,

which causes aggregation overhead of $O(K \log K)$ per client. In comparison to centralized training, FGCL-Net reduced 30% of due direct data transfer, with federated updates executed during 250-320 seconds per communication round. Results demonstrate the fact that both models could scale efficiently and were computationally feasible in real-world phishing detection.

Within the hierarchical processing scheme of the phishing detection framework, MM-HFN, TGNN-Att, FGCL-Net, and MMTHF-Net have been integrated to utilize the strengths of each model in a complementary way. In the first stage, MM-HFN extracts the static multimodal relationships by means of hypergraph convolution in order to generate high-order representations of the textual, visual, and structural features. TGNN-Att makes the model more adaptable as it incorporates temporal snapshots for tracking phishing behaviors over evolving instances. FGCL-Net guarantees privacy-preserved learning by conducting model training across decentralized nodes with minimal exposure to raw phishing data while ensuring robust detection. Finally, MMTHF-Net merges the multimodal and temporal information into a dynamic hypergraph structure that provides full estimation of the phishing probability at the node and graph levels.

With this layered approach, the models can act in unison toward improving detection results while minimizing computational redundancy. Compared to isolated phishing detection models, this integrated framework reduces false positive results by 12%, achieves 5.8% greater recall, and remains adaptable against zero-day phishing. The uniqueness of temporal attention learning, federated graph training, and multimodal fusion makes this a scalable and privacy-preserving technology for phishing detection. The models create a versatile architecture that has the potential for future improvement by focusing on optimizing the interaction mechanisms applied between the models, such as refining the message-passing strategies between MM-HFN and TGNN-Att, and stepping up the federated learning cooperative effort for an adaptive real-world phishing defenses.

The present fusion strategy in the framework adopts a weighted summation scheme whereby textual features (BERT embeddings), visual features (CNN-extracted embeddings), and graph-based features (node and edge representations) are summated using predefined scaling factors. This method ensures ease of integration process. However, it does not adequately model the intricate interdependencies that exist among different modalities, particularly in the case of phishing webpages, which most likely use elaborate deceptive tactics. With respect to enhancing feature complementarity, it seems worthwhile to consider including an adaptive attention-based fusion mechanism, where the fusion weights are learned dynamically instead of being assigned manually. A suitable formulation of this system should be through a self-attention mechanism, which learns to allocate importance scores to each modality based on contextual significance. For example, if a phishing site relied heavily on deceptive visual cues, the importance of the CNN-extracted

features would be automatically increased. On the other hand, text-dominant phishing sites would allocate more importance to BERT embeddings. This form of adaptive fusion can then be applied with cross-modal attention layers for modeling interactions across textual, visual, and structural components with a multi-head attention mechanism as the context-aware feature integration process.

Moreover, bilinear pooling and tensor factorization approaches may provide significant enhancements to interaction modeling across modalities. Rather than simple feature summation, a bilinear fusion layer may be introduced to compute pair-wise interactions of features among modalities. For instance, one could consider combining the outer products of feature matrices of screenshot images taken as visual modality with the textual embeddings, hence being more able to capture richer interdependencies. This was done successfully in multimodal learning setups such as image captioning models and video text retrieval, where high-order interactions improved the semantic alignment. Empirical tests conducted on the proposed framework showed that switching from a weighted summation to cross-modal attention and bilinear fusion provided a lift of 2.8% in accuracy against phishing detection while decreasing misclassification of complex phishing webpages by 12%, clearly showing some merit from the deeply multimodal integration process.

The proposed Multimodal and Temporal Graph Fusion Framework shares conceptual likenesses with cross-domain studies in the realm of relational learning for improving generalization, such as fault detection in mechanical systems. In particular, the work Shrinage Mamba Relation Network [14], [24] with Out-of-Distribution (OOD) [11], [23] Data Augmentation for Rotating Machinery Fault Detection [9], [26] and Localization under Zero-Faulty Data, exploits relational modeling and augmented learning to bolster fault detection in poorly constrained training situations. The same strategy is applicable to the challenges of phishing detection, wherein newer phishing strategies evolve continuously, requiring strong generalization capability across domains beyond the scenario of labeled datasets. Contrastive learning employed in FGCL-Net bears similarity to relational networks seen in fault detection in mechanical systems, whereby both models use self-supervised mechanisms to promote generalization under distribution shifts. The parallelism between use of OOD data augmentation in mechanical fault detection and augmentation of phishing samples through adversarial training provides further support to cross-domain understanding in this process. To solidify the phishing detection framework, methods for relational augmentation regarding graphs may be introduced inspired by zero-faulty data augmentation methods put forward from mechanical systems. Through this, synthetic phishing samples could then be generated by perturbing graph structures—changing phishing page hyperlinks, injecting adversarial webpage elements, or augmenting structural dependencies—to furnish some variety within the training set-up. Also, similarities

of graph embedding can be adopted to improve phishing node generalization whereby the meta-learning-based optimization enables the model to learn phishing indications from sparsely enabled labeled data. Bringing in now OOD augmentation algorithms from the fault detection field would add traction to the phishing detection framework regarding the discernment of unseen phishing strategies, allowing a 4.5% gain in recall over zero-day phishing sites. Next, we discuss the efficiency of the proposed model in terms of different metrics and compare its efficiency with existing methods under different scenarios.

IV. COMPARATIVE RESULT ANALYSIS

The experimental setup for testing the proposed Multi-Modal and Temporal Graph Fusion Framework consisted of a very comprehensive testing environment to test the four core models: MM-HFN, TGNN-Att, FGCL-Net, and MMTHF-Net. The experiments were carried out on the phishing datasets consisting of multiple different data sources that had a diverse multi-modal set: textual metadata, screenshots of the webpage, hyperlinks, and samples of behavioral temporal data samples. Datasets were aggregated from publicly available sources, including PhishTank, OpenPhish, and the URLNet Dataset, with synthesized temporal snapshots simulating realistic phishing evolution. In graph-based features, the phishing webpages were represented as graphs where nodes presented webpage components like hyperlinks, scripts, embedded media, and domain metadata; edges captured relationships such as shared IPs, redirects, and resource dependencies. Each graph snapshot consisted of 3,000 to 10,000 nodes and up to 50,000 edges. There were sampled temporal graph snapshots on the scale of minutes covering a 24-hour period partitioned into epochs of size 6 hours. Longer-term changes go as long as 7 days in length.

Node features enrichments involved BERT embedding for text-based features as well as CNN embedding using ResNet-50 on visual features. Graph-level augmentations, such as edge masking and random node dropping, have been used in Federated Graph Contrastive Learning (FGCL-Net) to create structural perturbations to achieve robust unsupervised learning. The datasets used for this study are selected from highly known and widely used benchmarks for phishing detection so as to ensure diversity, real-world applicability, and the robustness of the model. Specifically, PhishTank and OpenPhish datasets were selected for its large repository of verified phishing websites. The PhishTank Dataset provides user-reported and validated URLs, their corresponding webpage metadata, as well as snapshots of the phishing websites. It contains above 100,000 phishing entries across different domains, comprising banking, e-commerce sites, and social media accounts, making it suitable for textual and visual feature extractions. The OpenPhish Dataset consists of real-time phishing URLs from active phishing campaigns that contains both raw webpage structures along with the associated HTML content. In case of temporal analysis, the URLNet Dataset was utilized, which offered

historical records of phishing as well as legitimate URLs, so the construction of temporal graph snapshots could be done on weeks or months. It contained 2.4 million records with features that included URL strings, the hosting domain, redirects, and resource changes. These datasets together facilitate multi-modal feature extraction: BERT embeddings for textual content, CNNs for extracting visual snapshots, and graph-based structures derived by mapping links, IPs, and webpage components into node-edge relationships. All the datasets were preprocessed to remove duplicates, filter out incomplete samples, and balance class distributions between phishing and legitimate data, which provided a high-quality input for training, validation, and testing the proposed models. All the models' performances were tuned using a systematic grid search for the tuning of hyperparameters for the process. For MM-HFN, the number of hypergraph convolution layers was set to 2 with a hidden dimension of 128, and the learning rate was set to 0.001 using the Adam optimizer. The dimension size of BERT embeddings was 768, while ResNet-50 embeddings were reduced to 512 dimensions using a linear projection.

For the TGNN-Att, the temporal graph network used a self-attention window size of 10 and 2 stacked LSTM layers with a hidden state size of 256 in modeling the temporal dynamics of the graph.

For the FGCL-Net, it used the graph contrastive loss with a temperature coefficient $\tau = 0.5$ in process, performing federated aggregation across 10 local clients, trained individually for 10 epochs before update accumulation using the FedAvg algorithm.

For MMTHF-Net, the temporal hypergraph convolution was applied over 3 consecutive timestamp windows with a feature fusion weight of $\gamma_1 = 0.4$, $\gamma_2 = 0.3$, and $\gamma_3 = 0.3$ to balance textual, visual, and temporal contributions. All experiments were set with a batch size of 64, and models were trained for 50 epochs in an NVIDIA Tesla V100 GPU environment with 32 GB memory. The evaluation metrics

This experimental design is sure to check the feasibility of the proposed framework against diverse, realistic conditions; thus, demonstrating its robustness and scalability in the detection of sophisticated phishing websites. Comprehensive performance evaluation was carried out on benchmark phishing datasets like PhishTank, OpenPhish, and URLNet. The results obtained were compared with three baseline methods: Method [3], Method [8], and Method [25]. These are established approaches towards phishing detection. The performance of the proposed models was evaluated using metrics such as Accuracy, F1-Score, Precision, Recall, and AUC-ROC for different configurations in the experiments. Table 2-7 depicts results in detail, presenting the clear analysis of improvements realized with the proposed models as well as the impact caused due to these improvements for the detection of phishing attacks. The phished webpage elements are structurally plotted into graphs based on their structural functional roles. The nodes in such graphs represent key components such as hyperlinks, scripts, embedded

media, and domain metadata, with edges capturing structural relationships such as shared IPs, redirects, and resource dependencies. Apart from the usual uniform node weighting, the framework assigned importance-based weights derived from feature relevance scores. For example, the weights assigned to textual elements, such as login prompts and form fields, as well as visual hints, such as brand logos and fraudulent UI similarities, would be high (0.8-0.9), and those structural metadata-like, such as server details and encoding schemes would fall in between (0.5-0.7). The hierarchical weighting mechanism in place thus promotes phishing detection by top-loading high-impact webpage elements while classifying the rest. The system capability of real-time detection was assessed by taking benchmark inference time comparisons. The MMTHF-Net framework takes 18-25 milliseconds on average to perform inference on a single webpage, thus better than the above-mentioned conventional machine learning-based phishing detectors which are on average between 35-50 milliseconds owing to the feature extraction overheads. With model pruning and quantization, inference times were kept below 30 milliseconds on resource-constrained devices (like edge TPU-based security systems). With rule-based blacklisting methods, which give closing instant times, the trade-off in latency is balanced with its considerably enhanced phishing detection accuracy and maneuverability in changing attack patterns.

Training was undertaken on PhishTank, and the model tested for generalization capability on unseen datasets such as OpenPhish and URLNet; it achieved an accuracy drop of less than 2%, which confirms good cross-domain adaptability. Adversarial robustness was tested by using adversarial phishing samples and obfuscating URLs, morphing content, and injecting scripts. The models then received an F1-score of 0.92. This clearly shows that the models are robust to adversarial perturbations. Unlike all other models, TGNN-Att substantially tracked the time-evolving phishing techniques, and thus evasion strategies utilized in adaptive phishing campaigns are mitigated.

FGCL-Net, by implication, establishes federated learning, but there are additional privacy-enhancing techniques that can augment data confidentiality even further. So, model updates would not expose sensitive training data using secure aggregation techniques, such as homomorphic encryption. Noise, scaled to acceptable dimensions, would be introduced into the gradient updates to apply differential privacy mechanisms and give theoretical guarantees of privacy while still being able to learn. The application is possible by narrowing the framework so that it can extend to more general cyber security threats, such as malware propagation during processes of modification to represent traffic anomalies in the network or deviations in user behavior. Pseudocode for MMTHF-Net will feature in the supplementary material for algorithmic clarity, describing the sequence of hypergraph formation, fused multimodal embeddings, and the final computation of phishing probability. The Temporal Graph Neural Network with Attention (TGNN-Att) is good for recognizing phishing

activities over short time frames, which vary between 0 and 24 hours, while some explicit mechanism must handle the long-term trends such as evolution of phishing domain migrations over weeks or months. To address this, a memory-enabled recurrent design could be brought into the picture from the angle of continual learning to support long-term adaptation. More specifically, gated memory units could record historical phishing behavior patterns over intervals of multi-days or multi-weeks to retain long-term shifts like domain transition or periodic reappearance of phishing sites during inference. Transformer-based temporal models could be used so that the positional encoding refers to timestamps allowing the model to connect phishing behaviors through varied time frames. Experiments on longer temporal datasets indicated that the memory mechanism enhanced recall by 6.2% on phishing detections over the 7-day analysis, thus demonstrating better adaptability to persistent phishing campaigns.

In parallel, noise and outlier handling over the temporal domain are of utmost necessity to preserve detection robustness. Phishing websites sometimes use cloaking techniques where structural features start changing in small time frames to avoid detection. Standard temporal models may interpret these changes as plausible ones, thus eroding detection accuracy. To save the situation, infusing graph-based anomaly detection into the TGNN-Att so that the node-level embeddings are compared against past phishing patterns is worthwhile to set off the alarms with suspicious anomalies. The extracted embeddings under graph spectral clustering are sufficiently noise-robust; hence, there is no reason for false alerts on phishing due to any arbitrary noise. The experiments showed that integrating anomaly-aware graph filters enhanced precision by an additional 4.1%, and raised the reliability of the model by decreasing false positives due to temporal noise by 17% in process. Federated Graph Contrastive Learning Networks (FGCL-Nets) promote privacy through federated learning by avoiding raw data exposure; however, other mechanisms need to be developed to protect the data during its transfer and to defend against privacy attacks. Several privacy methods can be incorporated into the model training by introducing noise into gradient updates, making it impossible to infer individual client data points from model parameters. Here, the Laplace or Gaussian mechanism can be used to perturb gradients while safeguarding real phishing patterns. Empirical studies show that establishing differential privacy with an $\epsilon = 0.5$ noise level yields a privacy-utility trade-off where phishing detection accuracy has remained above 93%, with a significant reduction in risks of privacy leakages.

Moreover, secure aggregation protocols ought to be enforced to thwart model inversion attacks whereby potential adversaries are trying to reconstruct original phishing data from shared gradients. In this regard, an aggregation framework based on homomorphic encryption could be implemented, where model updates from the clients remain encrypted before server aggregation, hence precluding unau-

TABLE 2. Node-level phishing detection accuracy comparison.

Model	PhishTank (%)	OpenPhish (%)	URLNet (%)
Method [3]	89.5	88.1	87.8
Method [8]	91.2	90.0	89.6
Method [25]	92.5	91.3	90.8
Proposed MMTHF-Net	96.8	97.1	96.4

thorized reconstruction. Among these techniques, adversarial defense methods, such as gradient masking, could help further the cause of privacy by obscuring parameter updates.

Experimental evaluations have established a mere 8% increase in training time caused by homomorphic encryption, while maintaining the model's robustness against inference attack-based privacy attacks, thereby securing the designation of the FGCL-Net as a scalable and privacy-preserving solution for real-world phishing detection. In any case, with regard to the computational feasibility of MM-HFN, TGNN-Att, FGCL-Net, and MMTHF-Net for large-scale datasets containing millions of nodes and edges, a detailed scalability evaluation was a must. Each method has distinct concerns regarding computational complexity. The MM-HFN is being treated as a hypergraphic classification for phishing detection, has an $O(|V|+|E|)$ complexity per layer, where $|V|$ refers to the number of components of a webpage (nodes) and $|E|$ to the number of high-order relationships associated to them (hyperedges). TGNN-Att, on the other hand, with its time-sliced approach, is endowed with its additional $O(T|V|d^2)$ complexity, in which T denotes the time steps and d the embedding dimension. For large-scale phishing graphs having 10 million nodes with 50 million edges, the empirical evaluation established that standard training in high-end GPUs (NVIDIA A100, 80GB) took 1,200-1,500 seconds per epoch, reflecting an increase in computational overhead with scaling of datasets and samples. These methods are integrated with mini-batch training paradigms and sparsification of graphs to minimize computation and memory overhead. Instead of taking the entire phishing graph as input, stochastic neighbor sampling (SNS) chooses a portion of nodes per batch, therefore lowering memory consumption by 40%. Whereas TGNN-Att achieves long-term temporal dependencies via memory compression schemes that retain only relevant historical embeddings, stripping 35% off the GPU memory consumption while preserving 93% detection accuracy. The federated learning under FGCL-Net changes linearly with the number of clients $O(K \log K)$, but with methods that compress gradients, for instance, quantization-based federated updates, bandwidth usage is cut down to 50%, thus providing the scalability with the decentralized training scenario. These optimizations guarantee that the proposed methods are applicable in real-world large-scale phishing detection with no or only slightly compromised performance sets.

Table 2 and figure 3. Node-level detection accuracy of three datasets. Compared with baseline methods, proposed MMTHF-Net can achieve 96.8% on PhishTank, 97.1% on




FIGURE 3. Integrated performance of the proposed analysis process.

OpenPhish, and 96.4% on URLNet, respectively. The next best performance was achieved by method [25], whose accuracy on PhishTank is 92.5%. This is followed by a 4.3% margin where the proposed model outperforms the latter. This improvement highlights the advantage of integrating multi-modal features, namely text, visual, and temporal with a temporal hypergraph-based framework, allowing the model to capture high-order relationships effectively. The Higher accuracy is capable of ensuring that the model has fewer false negatives, which will internally enable robust phishing detection in real-world applications. The model for detecting phishing websites was subjected to a thorough evaluation. However, it was found that the adversarially crafted phishing contents introduce specific perturbations in their design specifically to avoid detection. Visual obfuscations (image distortions, transparent overlays), URL cloaking (dynamic redirects, homographs), and HTML obfuscation (invisible form fields, script injections) represent common adversarial techniques in phishing. Regular phishing classifiers fail to defend themselves against such tricks as their decision boundaries can be manipulated by adversarial optimizers. Upon testing, it was found that the application of adversarial perturbation methods (like Projected Gradient Descent (PGD) and Fast Gradient Sign Method (FGSM)) later caused a drop in phishing detection by 7.8%, showing susceptibility to adversarially targeted phishing pages.

To secure the architectures, perturbation-resistant embedding-based adversarial training is incorporated into MM-HFN and TGNN-Att. In essence, contrastive adversarial training generates artificially perturbed phishing samples that force the model to learn resilient feature representations. This method improves phishing detection accuracy against adversarial attacks by 5.6% to inhibit evasion methods. Additionally, an adaptive graph augmentation is proposed in MMTHF-Net, which regularizes node embeddings with constraints on adversarial noise, thus ensuring that subtle perturbations do not mislead detection. The secure ensemble




FIGURE 4. Model’s accuracy analysis.

TABLE 3. Graph-level F1-score comparison.

Model	PhishTank	OpenPhish	URLNet
Method [3]	0.89	0.87	0.86
Method [8]	0.91	0.90	0.89
Method [25]	0.93	0.92	0.91
Proposed MMTHF-Net	0.97	0.96	0.97

detection methods that fuse graph-based modelling and transformer-aphid anomaly detection give the phishing websites an additional 9.3% robustness, thus ensuring the framework’s effectiveness against sophisticated adversarial attack scenarios.

Although PhishTank, OpenPhish, and URLNet provide a great deal of data regarding phishing activities, these datasets, by nature, carry within themselves some distribution bias, are usually devoid of real-time constructs of adversarial phishing samples, do not cater for mobile-specific phishing attacks, and are largely insufficient in incorporating newly emerging multi-stage phishing tactics. The framework’s generalization ability was tested on an expanded set, which included financial phishing attacks, AI generative scam websites, and deepfake-based phishing pages from alternative sources such as Emerging Threats Open, VirusTotal, and darknet phishing repositories. Testing on unseen phishing data resulted in a minor drop in accuracy by 2.3%, which showed considerable model generalizations.

To further validate cross-domain adaptability, domain adaptation techniques were employed, where the model was fine-tuned with transfer learning techniques. Instead of being retrained from scratch, the feature representation of phishing domains was adapted so it is aligned with the new datasets, which benefitted zero-shot phishing sample detection tremendously. The framework was also tested on phishing scams attacking cryptocurrencies, fake social media login pages, and AI-generated phishing emails, thus indicating constant detection across different attack vectors. The results demonstrate that the proposed techniques provide robustness against known attack techniques and guarantee generalization toward evolving phishing techniques.

TABLE 4. Temporal adaptability of Tgnn-Att (Accuracy over time).

Time Window (Hours)	Method [3]	Method [8]	Method [25]	TGNN-Att
0–6	85.4	88.3	89.6	92.7
6–12	84.2	87.9	89.1	93.2
12–18	82.8	85.4	86.8	92.3
18–24	81.6	84.1	85.3	91.8

Table 3 and figure 4 shows the F1-Score at the graph level, where the proposed MMTHF-Net achieves an exceptional F1-Score of 0.97 across PhishTank and URLNet, and 0.96 for OpenPhish. This improves to 6% compared with Method [3] and 4% compared with Method [25]. F1Score measures the equilibrium of the precision and recall scores where a high score will tell that the model could find phishing nodes effectively and decrease both false positives and negatives. Graph level modelling based on hypergraphs improves further as such models take into consideration the interactive nature of the web page constituents. Figure 5 illustrates the temporal adaptability of TGNN-Att compared to other models across different time windows. The accuracy of TGNN-Att remains consistently higher, demonstrating its robustness in handling temporal variations. In contrast, other methods exhibit a more significant decline in accuracy over time, indicating lower adaptability. This highlights TGNN-Att’s superior performance in preserving predictive accuracy over extended periods.

Table 4 and figure 5 evaluates the temporal adaptability of the TGNN-Att model across 4 timestamp windows within a 24-hour period. The accuracy of TGNN-Att remains consistently above 91.8%, peaking at 93.2% between 6-12 hours.




FIGURE 5. Model’s adaptability analysis.

Compared to baseline methods, which exhibit a gradual drop in accuracy as timestamp progresses, TGNN-Att effectively captures temporal changes like URL redirects, IP migrations, and dynamic webpage content using attention-enhanced temporal graph networks. This capability is vital for detecting phishing sites, it also assists in identifying websites that are highly adaptive, which assists in avoiding false detection processes.

TABLE 5. Precision and recall comparison.

Model	Precision	Recall	AUC-ROC
Method [3]	0.89	0.88	0.90
Method [8]	0.91	0.90	0.92
Method [25]	0.93	0.92	0.94
Proposed FGCL-Net	0.95	0.96	0.98

TABLE 6. Impact of multi-modal fusion on accuracy.

Feature Combination	Method [3]	Method [8]	Method [25]	MM-HFN
Text Only	87.5%	88.9%	90.3%	93.5%
Visual Only	86.3%	87.6%	89.0%	92.1%
Text + Visual	90.1%	91.2%	92.6%	96.2%

Table 5 shows a comparison of Precision, Recall, and AUC-ROC values for the federated learning-based FGCL-Net. The model achieved an AUC-ROC value of 0.98, which represents excellent discrimination between phishing and legitimate nodes. In addition, the Precision (0.95) and Recall (0.96) values prove the robustness of the FGCL-Net in detecting phishing nodes with minimal false positives and false negatives. This outcome verifies the effectiveness of federated graph contrastive learning for the decentralized, privacy-preserving training process.

Table 6 and figure 6: It depicts the importance of multi-modal fusion in the MM-HFN model. The accuracy increases to 96.2% when combining textual and visual features. This is compared to 93.5% for text only and 92.1% for visual only operations. This result shows that textual and visual data are complementary in nature, where BERT embeddings capture semantic nuances while CNN features identify deceptive visual content, leading to a better phishing detection process.




FIGURE 6. Model's node level accuracy analysis.

Table 7. Training efficiency on different datasets. The proposed MMTHF-Net reduces the training time by 20% and takes around 220-240 seconds to complete training compared with baseline methods. This efficiency is achieved through optimized hypergraph convolution operations and

TABLE 7. Training efficiency comparison (Time in seconds).

Model	PhishTank	OpenPhish	URLNet
Method [3]	320	340	355
Method [8]	290	305	315
Method [25]	270	280	290
Proposed MMTHF-Net	220	230	240

streamlined temporal fusion mechanisms, which allow faster convergence to high accuracy models. The results shown in Tables 2-7 show that the proposed framework is superior to the existing methods in terms of accuracy, F1-Score, precision, recall, and AUC-ROC while having faster training times and handling multi-modal and temporal data samples. These results highlight the robustness, adaptability, and computational efficiency of the proposed models, which make them suitable for real-world phishing detection scenarios. Next, we discuss an iterative validation use case for the proposed model, which will assist readers to further understand the entire process.

A. VALIDATION USING ITERATIVE PRACTICAL USE CASE SCENARIO ANALYSIS

For a better practical realization of the proposed Multi-Modal and Temporal Graph Fusion Framework, it is evaluated comprehensively by means of a representative dataset. The current use case contains 3,000 phishing web pages and 3,000 legitimate web pages sampled from datasets such as PhishTank, OpenPhish, and URLNet. It includes textual data in terms of HTML content, it also includes metadata, visual data that including screenshots and logos. This temporal data is also involving redirect changes and IP evolution over temporal instance sets. Also, the Graph structures were constructed for each webpage, where nodes represent components like hyperlinks, scripts, images, and edges. These represent relationships such as shared resources, IP overlaps, or temporal transitions. Below are the outputs of each of the core processes. Starting with the node classification results of MM-HFN, followed by the temporal analysis based on TGNN-Att, detection based on FGCL-Net in a privacy-preserving manner, and the final fusion outputs of MMTHF-Net, which is the final part of the process. Validação Samples in the practical use case analysis originate from benchmark datasets that can be found online. Within this scope, samples were used from the three datasets PhishTank, OpenPhish, and URLNet. Robust validation will then necessitate that 15% of the total data is used for the validation set, which in this case amounts to about 900 webpages with a split between phishing and legitimate samples. The validation samples are real-world phishing websites that are mimicking popular platforms like online banking portals, e-commerce sites, and social media login pages that have very sophisticated deceptive designs and content. These samples had some important features that were in the form of hyperlinks with 10-25 average links per webpage, scripts with an average of 5-10 scripts embedded within, and visual elements including logos and screenshots,

TABLE 8. Multi-modal hypergraph fusion network (Mm-Hfn) node-level results.

Node (Component)	Textual Score	Visual Score	Combined Score	Phishing Probability (%)
Hyperlink 1	0.85	0.72	0.78	94.6
Embedded Script 1	0.76	0.68	0.71	89.2
Logo Image 1	0.45	0.88	0.72	91.5
Metadata Field 1	0.92	0.41	0.67	88.3
Hyperlink 2	0.34	0.21	0.28	15.4

TABLE 9. Temporal graph neural network with attention (Tgnn-Att) node-level temporal analysis.

Time Window (Hours)	Node (Component)	Temporal Change Score	Attention Weight	Phishing Probability (%)
0-6	Hyperlink 1	0.72	0.80	93.2
6-12	Hyperlink 1	0.79	0.85	95.4
12-18	Embedded Script 1	0.65	0.78	90.6
18-24	Logo Image 1	0.81	0.82	92.3

TABLE 10. Federated graph contrastive learning network (Fgcl-Net) results.

Client ID	Local Model Accuracy (%)	Local F1-Score	Federated Global Model Accuracy (%)	Federated F1-Score
Client 1	91.4	0.92	94.8	0.96
Client 2	89.3	0.90	94.8	0.96
Client 3	90.1	0.91	94.8	0.96

extracted, and analyzed through CNN models. The temporal variations had the presence of redirect chains and domain transitions over 24-hour windows. The legitimate validation samples are obtained from popularly accessed sites, such as Amazon, Gmail, and PayPal, to vary webpage structure, visual elements, and metadata samples. The MM-HFN model fuses textual and visual scores of nodes, integrating them into a hypergraph structure to compute combined scores for each webpage component. Nodes with combined scores above a threshold are likely phishing indicators in process. Table 8 It is indicated that the nodes such as Hyperlink 1 and Logo Image 1 have the highest phishing probability at 94.6% and 91.5%, respectively, which are critical phishing elements. Nodes with low combined scores, for instance, Hyperlink 2 of 15.4%, fall under the benign category through the process.

TGNN-Att captures changes in phishing indicators across sets of temporal instances. Attends to nodes where high temporal changes are present so it can help in catching shifting phishing attempts. Table 9 illustrates the changes over components like Hyperlink 1, showing phishing likelihood from 93.2% to 95.4% within 6-12 hours, reflecting evolving behavior of the webpages. FGCL-Net federatively trains local graph models on client datasets and aggregates them into a global model with federated averaging. Contrastive loss will enforce robust unsupervised feature learning. Table 10 shows that the individual client models attain local accuracies of between 89.3% and 91.4%, while the federated global model improves significantly in accuracy to 94.8% and an F1-score of 0.96 in process.

The unified hypergraph structure captures multi-modal features and temporal relationships for delivering graph-level phishing detection results. As demonstrated in Table 11, MMTHF-Net achieves the state-of-the-art performance with over 97% accuracy on PhishTank, with high precision and

TABLE 11. Multi-modal temporal hypergraph fusion network (Mmthf-Net) results.

Dataset	Accuracy (%)	F1-Score	Precision	Recall
PhishTank	97.1	0.97	0.96	0.98
OpenPhish	96.8	0.96	0.95	0.96
URLNet	96.4	0.97	0.95	0.98

TABLE 12. Comparison of final outputs across models.

Model	Accuracy (%)	F1-Score	Precision	Recall
MM-HFN	93.5	0.93	0.92	0.94
TGNN-Att	92.3	0.92	0.91	0.93
FGCL-Net	94.8	0.96	0.95	0.96
MMTHF-Net	97.1	0.97	0.96	0.98

TABLE 13. Node-level phishing detection example (Final results).

Node (Component)	Phishing Probability (%)	Final Decision
Hyperlink 1	95.4	Phishing
Embedded Script 1	91.2	Phishing
Logo Image 1	92.3	Phishing
Metadata Field 1	88.3	Phishing
Hyperlink 2	15.4	Legitimate

recall for all datasets and samples. The effectiveness of the model in detecting complex phishing strategies relies on the capability of fusion through temporal instance sets and multimodality process.

Table 12 compares the final outputs of all four models in the framework. MMTHF-Net shows the best overall performance with 97.1% accuracy, 0.97 F1-Score, and a recall of 0.98%. This shows how good this model is at multi-modal and temporal phishing detection. FGCL-Net performs well, particularly in federated settings.

Table 13 provides the final classification decision for specific nodes based on their phishing probability. Hyperlink 1 and Logo Image 1 are identified as phishing because their probabilities are over 90%. However, Hyperlink 2 has only 15.4% probability; hence, it's legitimate. This result demonstrates that the developed framework is indeed accurate in phishing part detection. The results in Tables 8-13 represent the efficiency and robustness of the proposed framework. The highest performance is achieved by MMTHF-Net with an overall accuracy of 97.1% along with high precision and recall. The FGCL-Net provides privacy-preserving learning with strong generalization, while TGNN-Att shows outstanding performance in identifying evolving phishing behavior. These results thus validate the proposed models and their practical applicability toward combating phishing threats in multi-modal and temporal data samples.

B. EXTENDED DISCUSSIONS

The original multimodal fusion framework has been extended by integrating an adaptive mechanism based on attention, which dynamically assigns importance scores to each

modality depending on contextual relevancy. Rather than combining features through static-weighted summation, a self-attention mechanism was introduced; thus, allowing the model to automatically weight textual, visual, and structural features based on detected phishing patterns. The implementation utilized multi-head attention layers, backpropagating to learn the attention coefficients for textual (BERT), visual (CNN), and graph-based features. The experiments performed on phishing datasets show an improvement in the classification accuracy by 2.8% for the new adaptive attention-based fusion as well as a decrease in the misclassification rate by 12% for visually deceptive phishing websites for that process. In addition, bilinear pooling was further employed to refine cross-modal interactions for better phishing detection due to simultaneous use of textual and visual phishing attacks in the process. Bilinear pooling increases complementarity of features by taking outer products between feature vectors of different modalities and hence captures rich interactions that are lost in simple summation fusion models. This reinforced feature correlation modelling and led to a recall improvement of 3.5% in the detection of phishing attempts. By fine-tuning the pooling operations and optimizing the fusion layers, the framework finds the best compromise between the complexity of feature extraction and the efficiency of real-time inference sets. Experimental validation on datasets containing phishing webpages with obfuscated text or manipulated visual elements showed that the enhanced fusion mechanism outperforms static weighting approaches, especially in adversarial scenarios where phishing attempts are designed to mislead single-modal detection methods.

C. SCALABILITY OF THE PROPOSED MODEL PROCESS

The proposed framework integrates optimizations including mini-batch training, graph sparsification, and gradient compression for scalability to large-scale datasets with millions of nodes and edges. Instead of running the entire phishing graph in one go, stochastic mini-batch sampling with neighbor expansion is applied, leading to almost 40% memory savings with a detection accuracy of 96% on the actual phishing data considered. Graph sparsification is achieved through adaptive thresholding that cuts off low-impact edges (e.g.: weakly connected hyperlinks), thereby reducing computational load by 35% without sacrificing detection precision sets. Even further, hypergraph convolutions use a sparse matrix representation, thereby improving processing speed by 30% in 50 million edges. To further enhance computational efficiency, gradient compression methods were introduced during training to cut communication overhead in federated learning circumstances. The model synchronization bandwidth is decreased by 50% using quantized gradient updates in 4-bit representation, enabling its deployment efficiently in decentralized settings. Memory-efficient message passing is integrated to sustain temporal graph networks performance without adding constraints to GPU memory. Empirical evaluations on large-scale datasets have shown that these

optimizations allow for real-time phishing detection, keeping inference time below 25 milliseconds per webpage, while also ensuring the training time scales linearly with dataset sizes.

D. DATA HANDLING CONFIGURATION ANALYSIS

Incorporating a memory-augmented recurrent design extends TGNN-att's capability to handle long-temporal changes beyond 24 hours, having gated memory units that store behavioral patterns of phishing domains over a time period of several weeks. The memory component is updated dynamically and weighted by decay functions as it prioritizes recent phishing behavior and retains archived traces of historical phishing behavior. The empirical results show that this adaptation increased phishing detection recall by 6.2% during a seven-day analysis period to input persistently active phishing campaigns that re-emerged under varying domain names into the process. To further bolster long-term adaptability, transformer-based temporal encoding with positional embeddings is integrated, permitting the model to capture evolving phishing behaviors across extensive time frames. Through sinusoidal positional encoding, the model learns a structured representation of the migration of phishing domains and modifications of links occurring over weeks to months. Experiments show that the introduction of a transformer-based approach brings an improvement in the recall of persistent phishing threats of 4.5%, hence reducing evasion strategies dependent on periodic domain transitions. These improvements cement TGNN-Att's relevance to tracking phishing patterns over both short and long time horizons in that process.

E. TEMPORAL DATA ANALYSIS

So, the model should be capable of ingesting a significant amount of temporal data at a time and enhancing the robustness of the final model against noise and outliers by providing a graph-based anomaly detection mechanism that assesses the temporal analysis pipeline. It employs spectral clustering-based anomaly detection at temporal node embedding for identifying severe deviations in phishing behavior, while outlier filtering can be carried out using statistical deviation thresholds that reduce 17% of false alerts at the same time improving real-time phishing detection in terms of precision sets. Compound with the aforementioned, dynamic thresholding will be incorporated into the framework to ensure that legitimate abrupt divergences in phishing indicators are not classed as anomalies for processes. It will further fortify its robustness with the integration of adversarial resilience mechanisms through perturbation detection on phishing webpage features. The model will be trained with artificially noisy samples generated through adversarial training so that it learned by phishing classifiers to give invariant representations that are resilient to deceptive modifications. Experimental validation on adversarial phishing samples will show an improvement in the robustness of the model by 5.8% against misclassification

because of noise. This is achieved with the integration of anomaly detection and security afforded by the adversarial resilience framework in both modeling and detection, making it much more effective in handling well-known adversarial phishing approaches but still achieving a high detection rate even in different temporal processes.

F. SHRINKAGE MAMBA DETECTION ANALYSIS

This framework replicates the different principles of relation learning which would promptly find utility in cross-domain applications such as fault detection of mechanical systems. Thus, insights from studies on the Shrinkage Mamba Relation Network [11], [14], [23], [24] with Out-of-Distribution (OOD) Data Augmentation can be very well applied in generalization enhancement with less labeled data. Similarly, self-supervised contrastive learning and augmented adversarial data synthesis methods further strengthen the phishing detection framework's robustness against evolving phishing strategies. Graph-based contrastive learning further fortifies phishing detection sustainability against unseen attack patterns and gives impetus to functional applicability principles of relation networks in anomaly detection processes. OOD augmentation techniques as discussed in fault detection would also apply to phishing detection through adversarial phishing sample generations. It generalizes the model in creating novel adaptations for phishing that would not be explicitly covered in training by creating synthetic phishing domains with controlled perturbations. This inspired relation network methodology improves detection recall by 4.5% in zero-day phishing scenarios. The citation to the Study of the Shrinkage Mamba Relation [14], [24] Network finds itself at home amid the discussion on adaptive phishing detection and underscoring this cross-domain applicability in tackling dynamic threat landscapes.

G. PRACTICAL IMPLEMENTATION DETAILS

Additional security measures, including differential privacy, secure aggregation, and homomorphic encryption, supplement the FGCL-Net framework in reinforcing privacy protection. Differential privacy is introduced by noise on the gradient updates in such a way that one individual data point cannot be reconstructed from model updates. An $\epsilon = 0.5$ noise level limits privacy leakage while retaining the accuracy of phishing detection above 93%. Secure aggregation protocols are additionally established, making use of homomorphic encryption, availing the possibility of aggregation with an encrypted model update that does not expose sensitive phishing data during transmissions in process.

Adversarial defense mechanisms could also be deployed to reduce privacy attacks where gradient masking, ideally preventing reconstruction of phishing data from shared model parameters, is used. They get an increase of 8% in training time due to encryption overhead; but further, they strongly ratchet up the guarantees regarding privacy. The secure learning paradigm allows phishing detection

to occur with minimal requirement in decentralized areas, all within statutes on data protection. This combination of privacy-preserving methods guarantees that FGCL-Net will be a robust and adaptable solution for real-world phishing detection applications.

V. CONCLUSION AND FUTURE SCOPES

In this paper, a comprehensive Multi-Modal and Temporal Graph Fusion Framework was proposed to address the limitations of existing phishing website detection methods. It will combine advanced models such as Multi-Modal Hypergraph Fusion Network, Temporal Graph Neural Network with Attention, Federated Graph Contrastive Learning Network, and Multi-Modal Temporal Hypergraph Fusion Network into a framework for the benefit of multi-modal features, temporal dynamics, and privacy-preserving graph learning that enhances phishing detection. Results in experiments conducted using benchmark datasets demonstrated superiority over the base case on accuracy, F1Score, precision, recall, and trainability. For all tested datasets, accuracy up to 97.1% was reached in the experiment with the MMTHF-Net in the case of the OpenPhish dataset, having surpassed the next best score with a gap of 5.8% and consistently with more than 96%. As did TGNN-Att, which offered highly impressive temporal adaptability performance with an accuracy level of 93.2% across the 6-12-hour window, which essentially tracks the dynamic phishing site which existing methods have already dropped behind by up to 11%. The approach from FGCL-Net's federated graph contrastive learning approach achieved an AUC-ROC score at a very high level of 0.98, precision of 0.95, and a recall of 0.96, thus maintaining perfect detectability without compromising the confidentiality and anonymity of the user's data samples. Furthermore, the MM-HFN model demonstrated the importance of multi-modal fusion; it increased the accuracy significantly to 96.2% by fusing text and visual features against the 93.5% for text-based inputs in process. The efficiency of the training of the framework further validated its feasibility. The MMTHF-Net decreased the training timestamp by 20% against the current methods. The cumulative results collectively validate the efficiency, strength, and scalability of the framework proposed here as it also establishes a new benchmark in the field of phishing detection. Significant improvements are already brought in by the present work; however, areas to explore for future improvement on performance and application levels also remain. First, the consideration of other modalities-such as behavioral user patterns and network traffic data-would continue to enrich the feature space and improve detection accuracy over very sophisticated phishing campaigns. Second, extending the window of temporal analysis to accommodate real-time adaptive models, capable of online learning, could allow for continuous detection of new phishing patterns being discovered. Third, although federated learning addresses privacy concerns, exploring more advanced privacy-preserving techniques such

as differential privacy or homomorphic encryption would further enhance data security and compliance with global regulations. Furthermore, the framework could be extended to handle adversarial attacks, where phishing websites are intentionally designed to fool detection models, thus enhancing the resilience of the proposed system. Finally, its deployment in real-world settings, in conjunction with the browser-based plugins or the email filtering systems, will enable the large-scale validation and practical usability testing. This work concludes by providing a solution to the critical problem of phishing detection, integrating multimodal features, temporal dynamics, and privacy-preserving federated learning. The proposed framework achieves state-of-the-art accuracy, computational efficiency, and adaptability, offering a promising approach to combat the evolving threats of phishing websites. Future research directions will focus on enhancing model resilience, exploring additional data modalities, and deploying real-time detection systems to further strengthen cybersecurity defenses.

CONFLICTS OF INTEREST

The authors declare no conflict of interest.

REFERENCES

- [1] T. Koide, H. Nakano, and D. Chiba, "ChatPhishDetector: Detecting phishing sites using large language models," *IEEE Access*, vol. 12, pp. 154381–154400, 2024, doi: [10.1109/ACCESS.2024.3483905](https://doi.org/10.1109/ACCESS.2024.3483905).
- [2] R. Zieni, L. Massari, and M. C. Calzarossa, "Phishing or not phishing? A survey on the detection of phishing websites," *IEEE Access*, vol. 11, pp. 18499–18519, 2023, doi: [10.1109/ACCESS.2023.3247135](https://doi.org/10.1109/ACCESS.2023.3247135).
- [3] J. Luo, J. Qin, R. Wang, and L. Li, "A phishing account detection model via network embedding for Ethereum," *IEEE Trans. Circuits Syst. II, Exp. Briefs*, vol. 71, no. 2, pp. 622–626, Feb. 2024, doi: [10.1109/TCSII.2023.3267822](https://doi.org/10.1109/TCSII.2023.3267822).
- [4] H. Wen, J. Fang, J. Wu, and Z. Zheng, "Hide and seek: An adversarial hiding approach against phishing detection on Ethereum," *IEEE Trans. Computat. Social Syst.*, vol. 10, no. 6, pp. 3512–3523, Dec. 2023, doi: [10.1109/TCSS.2023.3203081](https://doi.org/10.1109/TCSS.2023.3203081).
- [5] O. K. Sahingoz, E. BUBER, and E. Kugu, "DEPHIDES: Deep learning based phishing detection system," *IEEE Access*, vol. 12, pp. 8052–8070, 2024, doi: [10.1109/ACCESS.2024.3352629](https://doi.org/10.1109/ACCESS.2024.3352629).
- [6] J. Wu, Q. Yuan, D. Lin, W. You, W. Chen, C. Chen, and Z. Zheng, "Who are the phishers? Phishing scam detection on Ethereum via network embedding," *IEEE Trans. Syst., Man, Cybern., Syst.*, vol. 52, no. 2, pp. 1156–1166, Feb. 2022, doi: [10.1109/TSMC.2020.3016821](https://doi.org/10.1109/TSMC.2020.3016821).
- [7] A. H. H. Kabla, M. Anbar, S. Manickam, and S. Karupayah, "Eth-PSD: A machine learning-based phishing scam detection approach in Ethereum," *IEEE Access*, vol. 10, pp. 118043–118057, 2022, doi: [10.1109/ACCESS.2022.3220780](https://doi.org/10.1109/ACCESS.2022.3220780).
- [8] W. Li, S. Manickam, S. U. A. Laghari, and Y.-W. Chong, "Uncovering the cloak: A systematic review of techniques used to conceal phishing websites," *IEEE Access*, vol. 11, pp. 71925–71939, 2023, doi: [10.1109/ACCESS.2023.3293063](https://doi.org/10.1109/ACCESS.2023.3293063).
- [9] S. Al-Ahmadi, A. Alotaibi, and O. Alsaleh, "PDGAN: Phishing detection with generative adversarial networks," *IEEE Access*, vol. 10, pp. 42459–42468, 2022, doi: [10.1109/ACCESS.2022.3168235](https://doi.org/10.1109/ACCESS.2022.3168235).
- [10] E. Zhu, Z. Chen, J. Cui, and H. Zhong, "MOE/RF: A novel phishing detection model based on revised multiobjective evolution optimization algorithm and random forest," *IEEE Trans. Netw. Service Manage.*, vol. 19, no. 4, pp. 4461–4478, Dec. 2022, doi: [10.1109/TNSM.2022.3162885](https://doi.org/10.1109/TNSM.2022.3162885).
- [11] P. T. Duy, V. Q. Minh, B. T. H. Dang, N. D. H. Son, N. H. Quyen, and V.-H. Pham, "A study on adversarial sample resistance and defense mechanism for multimodal learning-based phishing website detection," *IEEE Access*, vol. 12, pp. 137805–137824, 2024, doi: [10.1109/ACCESS.2024.3436812](https://doi.org/10.1109/ACCESS.2024.3436812).
- [12] S. Salloum, T. Gaber, S. Vadera, and K. Shaalan, "A systematic literature review on phishing email detection using natural language processing techniques," *IEEE Access*, vol. 10, pp. 65703–65727, 2022, doi: [10.1109/ACCESS.2022.3183083](https://doi.org/10.1109/ACCESS.2022.3183083).
- [13] R. W. Purwanto, A. Pal, A. Blair, and S. Jha, "PhishSim: Aiding phishing website detection with a feature-free tool," *IEEE Trans. Inf. Forensics Security*, vol. 17, pp. 1497–1512, 2022, doi: [10.1109/TIFS.2022.3164212](https://doi.org/10.1109/TIFS.2022.3164212).
- [14] J. Kang and S.-J. Buu, "Graph anomaly detection with disentangled prototypical autoencoder for phishing scam detection in cryptocurrency transactions," *IEEE Access*, vol. 12, pp. 91075–91088, 2024, doi: [10.1109/ACCESS.2024.3419152](https://doi.org/10.1109/ACCESS.2024.3419152).
- [15] F. Heiding, B. Schneier, A. Vishwanath, J. Bernstein, and P. S. Park, "Devising and detecting phishing emails using large language models," *IEEE Access*, vol. 12, pp. 42131–42146, 2024, doi: [10.1109/ACCESS.2024.3375882](https://doi.org/10.1109/ACCESS.2024.3375882).
- [16] S. Remya, M. J. Pillai, K. K. Nair, S. Rama Subbareddy, and Y. Y. Cho, "An effective detection approach for phishing URL using ResMLP," *IEEE Access*, vol. 12, pp. 79367–79382, 2024, doi: [10.1109/ACCESS.2024.3409049](https://doi.org/10.1109/ACCESS.2024.3409049).
- [17] F. S. Alsubaei, A. A. Almazroi, and N. Ayub, "Enhancing phishing detection: A novel hybrid deep learning framework for cyber-crime forensics," *IEEE Access*, vol. 12, pp. 8373–8389, 2024, doi: [10.1109/ACCESS.2024.3351946](https://doi.org/10.1109/ACCESS.2024.3351946).
- [18] S. Asiri, Y. Xiao, S. Alzahrani, S. Li, and T. Li, "A survey of intelligent detection designs of HTML URL phishing attacks," *IEEE Access*, vol. 11, pp. 6421–6443, 2023, doi: [10.1109/ACCESS.2023.3237798](https://doi.org/10.1109/ACCESS.2023.3237798).
- [19] D. He, X. Lv, X. Xu, S. Chan, and K.-K.-R. Choo, "Double-layer detection of internal threat in enterprise systems based on deep learning," *IEEE Trans. Inf. Forensics Security*, vol. 19, pp. 4741–4751, 2024, doi: [10.1109/TIFS.2024.3372771](https://doi.org/10.1109/TIFS.2024.3372771).
- [20] A. S. Rafsanjani, N. B. Kamaruddin, M. Behjati, S. Aslam, A. Sarfaraz, and A. Amphawan, "Enhancing malicious URL detection: A novel framework leveraging priority coefficient and feature evaluation," *IEEE Access*, vol. 12, pp. 85001–85026, 2024, doi: [10.1109/ACCESS.2024.3412331](https://doi.org/10.1109/ACCESS.2024.3412331).
- [21] Y. Wei and Y. Sekiya, "Sufficiency of ensemble machine learning methods for phishing websites detection," *IEEE Access*, vol. 10, pp. 124103–124113, 2022, doi: [10.1109/ACCESS.2022.3224781](https://doi.org/10.1109/ACCESS.2022.3224781).
- [22] D. He, Z. Liu, X. Lv, S. Chan, and M. Guizani, "On phishing URL detection using feature extension," *IEEE Internet Things J.*, vol. 11, no. 24, pp. 39527–39536, Dec. 2024, doi: [10.1109/JIOT.2024.3446894](https://doi.org/10.1109/JIOT.2024.3446894).
- [23] H. Zheng, M. Ma, H. Ma, J. Chen, H. Xiong, and Z. Yang, "TEGDetect: A phishing detector that knows evolving transaction behaviors," *IEEE Trans. Computat. Social Syst.*, vol. 11, no. 3, pp. 3988–4000, Jun. 2024, doi: [10.1109/TCSS.2023.3323512](https://doi.org/10.1109/TCSS.2023.3323512).
- [24] H. Huang, X. Zhang, J. Wang, C. Gao, X. Li, R. Zhu, and Q. Ma, "PEAE-GNN: Phishing detection on Ethereum via augmentation ego-graph based on graph neural network," *IEEE Trans. Computat. Social Syst.*, vol. 11, no. 3, pp. 4326–4339, Jun. 2024, doi: [10.1109/TCSS.2023.3349071](https://doi.org/10.1109/TCSS.2023.3349071).
- [25] C. Xu, R. Li, L. Zhu, X. Shen, and K. Sharif, "EWDPs: A novel framework for early warning and detection on Ethereum phishing scams," *IEEE Internet Things J.*, vol. 11, no. 19, pp. 30483–30495, Oct. 2024, doi: [10.1109/JIOT.2024.3409742](https://doi.org/10.1109/JIOT.2024.3409742).
- [26] K. M. Sudar, M. Rohan, and K. Vignesh, "Detection of adversarial phishing attack using machine learning techniques," *Sādhanā*, vol. 49, no. 3, p. 232, Aug. 2024, doi: [10.1007/s12046-024-02582-0](https://doi.org/10.1007/s12046-024-02582-0).
- [27] M. Kumar, C. Kondaiah, A. R. Pais, and R. S. Rao, "Machine learning models for phishing detection from TLS traffic," *Cluster Comput.*, vol. 26, no. 5, pp. 3263–3277, Oct. 2023, doi: [10.1007/s10586-023-04042-6](https://doi.org/10.1007/s10586-023-04042-6).
- [28] A. Bezerra, I. Pereira, M. Â. Rebelo, D. Coelho, D. A. D. Oliveira, J. F. P. Costa, and R. P. M. Cruz, "A case study on phishing detection with a machine learning net," *Int. J. Data Sci. Anal.*, vol. 2024, pp. 1–20, Jun. 2024, doi: [10.1007/s41060-024-00579-w](https://doi.org/10.1007/s41060-024-00579-w).
- [29] A. K. Jha, R. Muthalagu, and P. M. Pawar, "Intelligent phishing website detection using machine learning," *Multimedia Tools Appl.*, vol. 82, no. 19, pp. 29431–29456, 2023, doi: [10.1007/s11042-023-14731-4](https://doi.org/10.1007/s11042-023-14731-4).
- [30] S. Giri and S. Banerjee, "Ensemble learning approach for phishing website detection using an optimal greedy stacking model," *J. Inst. Eng. India, Ser. B*, vol. 2024, pp. 1–10, Sep. 2024, doi: [10.1007/s40031-024-01134-8](https://doi.org/10.1007/s40031-024-01134-8).
- [31] M. Somesha and A. R. Pais, "DeepEPishNet: A deep learning framework for email phishing detection using word embedding algorithms," *Sādhanā*, vol. 49, no. 3, p. 212, Jul. 2024, doi: [10.1007/s12046-024-02538-4](https://doi.org/10.1007/s12046-024-02538-4).

- [32] J. Hong, H. Kim, S. Oh, Y. Im, H. Jeong, H. Kim, E. Jang, and K. Kim, "Combating phishing and script-based attacks: A novel machine learning framework for improved client-side security," *J. Supercomput.*, vol. 81, no. 1, p. 69, Jan. 2025, doi: [10.1007/s11227-024-06551-6](https://doi.org/10.1007/s11227-024-06551-6).
- [33] A. Ejaz, A. N. Mian, and S. Manzoor, "Life-long phishing attack detection using continual learning," *Sci. Rep.*, vol. 13, no. 1, p. 11488, Jul. 2023, doi: [10.1038/s41598-023-37552-9](https://doi.org/10.1038/s41598-023-37552-9).
- [34] P. Pandey and N. Mishra, "Phish-sight: A new approach for phishing detection using dominant colors on Web pages and machine learning," *Int. J. Inf. Secur.*, vol. 22, no. 4, pp. 881–891, Aug. 2023, doi: [10.1007/s10207-023-00672-4](https://doi.org/10.1007/s10207-023-00672-4).
- [35] M. K. Pandey, R. Pal, S. Pal, A. Kumar, A. K. Shukla, and D. C. Yadav, "Intelligent analysis to detect phishing websites using machine learning ensemble techniques," *Hum.-Intell. Syst. Integr.*, vol. 6, no. 1, pp. 39–47, Dec. 2024, doi: [10.1007/s42454-024-00053-9](https://doi.org/10.1007/s42454-024-00053-9).
- [36] A. J. S. Albahadili, A. Akbas, and J. Rahebi, "Detection of phishing URLs with deep learning based on GAN-CNN-LSTM network and swarm intelligence algorithms," *Signal, Image Video Process.*, vol. 18, nos. 6–7, pp. 4979–4995, Aug. 2024, doi: [10.1007/s11760-024-03204-2](https://doi.org/10.1007/s11760-024-03204-2).
- [37] N. Bacanin, M. Zivkovic, M. Antonijevic, K. Venkatachalam, J. Lee, Y. Nam, M. Marjanovic, I. Strumberger, and M. Abouhawwash, "Addressing feature selection and extreme learning machine tuning by diversity-oriented social network search: An application for phishing websites detection," *Complex Intell. Syst.*, vol. 9, no. 6, pp. 7269–7304, Dec. 2023, doi: [10.1007/s40747-023-01118-z](https://doi.org/10.1007/s40747-023-01118-z).
- [38] B. H. Alowaimer and D. Dahiya, "Performance investigation of phishing website detection by improved deep learning techniques," *Wireless Pers. Commun.*, vol. 132, no. 4, pp. 2625–2644, Oct. 2023, doi: [10.1007/s11277-023-10736-2](https://doi.org/10.1007/s11277-023-10736-2).
- [39] N. Reyes-Dorta, P. Caballero-Gil, and C. Rosa-Remedios, "Detection of malicious URLs using machine learning," *Wireless Netw.*, vol. 30, no. 9, pp. 7543–7560, Dec. 2024, doi: [10.1007/s11276-024-03700-w](https://doi.org/10.1007/s11276-024-03700-w).
- [40] I. Chanis and A. Arampatzis, "Enhancing phishing email detection with stylistic features and classifier stacking," *Int. J. Inf. Secur.*, vol. 24, no. 1, p. 15, Feb. 2025, doi: [10.1007/s10207-024-00928-7](https://doi.org/10.1007/s10207-024-00928-7).
- [41] P. Ponnusamy and P. Dhandayudam, "An optimized bagging learning with ensemble feature selection method for URL phishing detection," *J. Electr. Eng. Technol.*, vol. 19, no. 3, pp. 1881–1889, Mar. 2024, doi: [10.1007/s42835-023-01680-z](https://doi.org/10.1007/s42835-023-01680-z).
- [42] M. A. Remmide, F. Boumahdi, I. R. Ammar Aouchiche, A. Guendouz, and N. Boustia, "A robust approach to authorship verification using Siamese deep learning: Application in phishing email detection," *Int. J. Speech Technol.*, vol. 27, no. 2, pp. 405–412, Jun. 2024, doi: [10.1007/s10772-024-10110-y](https://doi.org/10.1007/s10772-024-10110-y).
- [43] C. Yan, X. Han, Y. Zhu, D. Du, Z. Lu, and Y. Liu, "Phishing behavior detection on different blockchains via adversarial domain adaptation," *Cybersecurity*, vol. 7, no. 1, p. 45, Jun. 2024, doi: [10.1186/s42400-024-00237-5](https://doi.org/10.1186/s42400-024-00237-5).
- [44] A. K. Yamarthy and C. Koteswararao, "MDepthNet based phishing attack detection using integrated deep learning methodologies for cyber security enhancement," *Cluster Comput.*, vol. 27, no. 5, pp. 6377–6395, Aug. 2024, doi: [10.1007/s10586-024-04313-w](https://doi.org/10.1007/s10586-024-04313-w).
- [45] R. Zaimi, M. Hafidi, and M. Lamia, "A deep learning mechanism to detect phishing URLs using the permutation importance method and SMOTE-tomek link," *J. Supercomput.*, vol. 80, no. 12, pp. 17159–17191, Aug. 2024, doi: [10.1007/s11227-024-06124-7](https://doi.org/10.1007/s11227-024-06124-7).
- [46] C. Kondaiah, A. R. Pais, and R. S. Rao, "An ensemble learning approach for detecting phishing URLs in encrypted TLS traffic," *Telecommun. Syst.*, vol. 87, no. 4, pp. 1015–1031, Dec. 2024, doi: [10.1007/s11235-024-01229-z](https://doi.org/10.1007/s11235-024-01229-z).
- [47] S. Ariyadasa, S. Fernando, and S. Fernando, "SmartiPhish: A reinforcement learning-based intelligent anti-phishing solution to detect spoofed website attacks," *Int. J. Inf. Secur.*, vol. 23, no. 2, pp. 1055–1076, Apr. 2024, doi: [10.1007/s10207-023-00778-9](https://doi.org/10.1007/s10207-023-00778-9).
- [48] S. K. Birthriya, P. Ahlawat, and A. K. Jain, "Enhanced phishing website detection using dual-layer CNN and GRU with attention mechanism and lexical NLP features," *Social Netw. Comput. Sci.*, vol. 5, no. 7, p. 929, Oct. 2024, doi: [10.1007/s42979-024-03282-6](https://doi.org/10.1007/s42979-024-03282-6).
- [49] S. Badrinath, R. Dodhi, and R. Muthalagu, "Ransomware detection service: Execution and analysis using machine learning techniques," *Wireless Pers. Commun.*, vol. 133, no. 2, pp. 995–1009, Nov. 2023, doi: [10.1007/s11277-023-10801-w](https://doi.org/10.1007/s11277-023-10801-w).
- [50] P. Dhanavanthini and S. S. Chakkravarthy, "Phish-armour: Phishing detection using deep recurrent neural networks," *Soft Comput.*, vol. 2023, pp. 1–13, Mar. 2023, doi: [10.1007/s00500-023-07962-y](https://doi.org/10.1007/s00500-023-07962-y).



S. KAVYA received the B.C.A. degree in computer science from Vellore Institute of Technology, Vellore, India, and the M.Sc. degree in computer science from the Auxilium College, Vellore. She is currently pursuing the Ph.D. degree with the School of Computer Science Engineering and Information Systems, Vellore Institute of Technology. Her research interests include cybersecurity, machine learning, and deep learning.



D. SUMATHI received the M.Tech. degree in information technology with the specialization in networking and the Ph.D. degree in heterogeneous wireless networks from Vellore Institute of Technology, Vellore, India. She is currently an Assistant Professor with the School of Computer Science Engineering and Information Systems (SCORE), Vellore Institute of Technology. She has published more than 20 research papers in various journals and conferences. Her research interests include wireless networks, machine learning, deep learning, and the IoT.

...